

MSCI 541 – Search Engines

Homework 4

Professor Mark Smucker

Matthew Erxleben

ID: 20889980

Date: November 21st, 2023

Problem 1)

Two documents that have the same BM25 retrieval score may differ in terms of relevance. This is due to the definition of document relevance. A document is deemed relevant due to the user of the query's own discretion. This is the only way that we can define relevance, as whether the document is relevant or not is only up to the user. Therefore, even if BM25 gives two documents the same score, one user may find one document relevant but not the other.


Problem 2)

2a)

In the equation for document term weights, the tf component contains a variable f_{jk} to represent the number of occurrences of term k in the document i . We use the log of this instead of just the variable in the formula. This is typically with the log base e , to be the natural logarithm. The reason for this to indicate the importance of the term k in the document i , versus just the raw number of how many times it occurs. The natural logarithm captures the exponential growth. If only the frequency was used, documents with more frequency would automatically be worth much more in the equation, which would out scale other documents with less words. The more important metric is indicating importance for a terms weight in a document, therefore that is why the log is used.

2b)

In the BM25 formula, the term frequency section uses a similar methodology. The frequency of term k in doc i is divided by k plus the frequency again. K is defined as a length normalizing coefficient, and is calculated using the doc length of doc i and the average doc lengths in the collection, as well as two tuning parameters b and $k1$. B is from 0 to 1, and controls the amount of length normalization (0.75 is typically used). $K1$ controls the saturation for term frequency in the doc (1.2 is typically used). This is similar to the previous implementation in the document term weights formula because it scales the frequency of terms appearing in the doc with how long the document is, and how long the average document is in the collection.



As we can see, the term frequency weight versus the frequency of the term is seen in the above chart, with corresponding document lengths. The term frequency weight is worth less when the document is greater in length.

Problem 3)

Typically, the frequency of terms in a query are ignored when doing BM25. However, we can use the query term frequency in the BM25 algorithm if wanted. We can do this by multiplying the terms contribution to the sum by the number of times it appears in the query. For example, if the query is “The brown Cleveland brown is brown in a brown town” and mentions the term brown four times, this would be valued differently than the query “The Cleveland brown is in a town” which mentions the term brown only once. Currently in our BM25 formulation, these queries are valued the same.

Problem 4)

The first step towards handling these words is would be stemming, if that is in the design of the algorithm. Stemming allows for words to be broken down into their basic form, and therefore increasing the chance that it may be seen in the document collection. For example, none of the documents may contain the term “eating”, but after stemming this word will become “eat”. This word may now be contained in a document that is in the collection; therefore, we consider that word as the query term instead. If stemming is not including in the algorithms design or still does not exist in the document collection, these words are typically handled by removing the words

from the query. This allows for certain algorithms to not get hurt from the term being included in the query terms, such as BooleanAND (which relies on all terms appearing at least once in each document)

Problem 5)

We never store this matrix as it is way to big. Instead, we use the count of terms from our inverted index to computer weights on the fly.

Problem 6)

Gary's has made many mistakes in his implementation of BM25.

The first mistake he made is converting all non-letters to spaces. This can cause words to break into words and remove the original term from the term list. For example, u.s.a will become three separate terms "u", "s", and "a" rather than becoming one term "usa". Gary would need to also include the full term "usa" for better results. Another mistake that Gary makes is not retaining the original tokens after stemming. It is important to test how the BM25 algorithm is with the original tokens, and then with the porter stemming to see what types of queries do better with the original tokens instead.

Problem 7)

Installation Requirements:

1. Please make sure Python is installed on your computer before running the program.
2. Clone the repository on your device by entering this into your terminal: `git clone https://github.com/UWaterloo-MSCI-541/msci-541-f23-hw4-matterxleben.git`

Running the Programs:

To run these programs, please navigate to where you cloned the repository and open the working directory `.../msci-541-f23-hw3-matterxleben`

IndexEngine:

This program accepts two command line arguments:

1. a path to the latimes.gz file
2. a path to a directory where the documents, metadata, and term files will be stored.

For example, you would run IndexEngine from the command prompt / terminal / shell as:

```
python IndexEngine.py C:/Users/matth/OneDrive/Desktop/University/3B/MSCI541-Search-Engines/HW4/raw-data/latimes.gz C:/Users/matth/OneDrive/Desktop/University/3B/MSCI541-Search-Engines/HW4/store
```

ResultsEvaluation.py:

The program accepts 2 command line arguments: the directory location of your results file, the directory location of your qrels file:

For example, you would run ResultsEvaluation.py from the command prompt / terminal / shell as:

```
python ResultsEvaluation.py C:/Users/matth/OneDrive/Desktop/University/3B/MSCI541-Search-Engines/HW4/hw4-files-2023/results/C:/Users/matth/OneDrive/Desktop/University/3B/MSCI541-Search-Engines/HW4/hw4-files-2023/qrels/LA-only.trec8-401.450.minus416-423-437-444-447.txt
```

BM25:

The program accepts three command line arguments: the directory location of your index, the queries file, and the name of a file to store your output

1. a path to the location of your index, created by IndexEngine
2. the queries file
3. the name of a file to store your output

For example, you would run BM25 from the command prompt / terminal / shell as:

```
python BM25.py C:/Users/matth/OneDrive/Desktop/University/3B/MSCI541-Search-Engines/HW4/store/term_files queries.txt hw4-bm25-baseline-merxlebe.txt
```

For this program, the queries file (search topics file) is in the /topics files folder. In the code for the program, the path is hardcoded as:

C:/Users/matth/OneDrive/Desktop/University/3B/MSCI541-Search-Engines/HW4/msci-541-f23-hw4-matterxleben/topics files/" + queries_file

Please change this to your own computers path to the msci-541-f23-hw4-matterxleben/topics files/ folder to obtain the queries (search topics)!

7b)

The results of the baseline BM25 retrieval can be found at hw4-bm25-baseline-merxlebe.txt

results >	hw4-bm25-baseline-merxlebe.txt					
1	401	Q0	LA101790-0075	1	6.074748903361086	merxlebeAND
2	401	Q0	LA021890-0100	2	6.032953328250141	merxlebeAND
3	401	Q0	LA050690-0109	3	5.6400852726916915	merxlebeAND
4	401	Q0	LA100890-0131	4	5.630885051378108	merxlebeAND
5	401	Q0	LA040789-0015	5	5.5405520405981825	merxlebeAND
6	401	Q0	LA021490-0049	6	5.5367825615584305	merxlebeAND
7	401	Q0	LA031590-0102	7	5.4670164670752435	merxlebeAND
8	401	Q0	LA071890-0073	8	5.42381110282183	merxlebeAND
9	401	Q0	LA040590-0157	9	5.383676256052823	merxlebeAND
10	401	Q0	LA050790-0042	10	5.379201506182712	merxlebeAND
11	401	Q0	LA082789-0152	11	5.367573570786622	merxlebeAND
12	401	Q0	LA050789-0068	12	5.332315133499008	merxlebeAND
13	401	Q0	LA111089-0188	13	5.32344019428111	merxlebeAND
14	401	Q0	LA021100-0168	14	5.307385769457557	merxlebeAND
15	401	Q0	LA060890-0011	15	5.299247923623829	merxlebeAND
16	401	Q0	LA020789-0133	16	5.229213772444686	merxlebeAND
17	401	Q0	LA030990-0189	17	5.222206946868615	merxlebeAND
18	401	Q0	LA050990-0043	18	5.193688047281026	merxlebeAND
19	401	Q0	LA062290-0172	19	5.183174087943154	merxlebeAND
20	401	Q0	LA090389-0171	20	5.162111967570969	merxlebeAND
21	401	Q0	LA050390-0176	21	5.136880574677754	merxlebeAND
22	401	Q0	LA043090-0043	22	5.124076675460465	merxlebeAND
23	401	Q0	LA112189-0066	23	5.122105045543815	merxlebeAND
24	401	Q0	LA091490-0080	24	5.120933739822092	merxlebeAND
25	401	Q0	LA061290-0074	25	5.111363185757659	merxlebeAND
26	401	Q0	LA021590-0209	26	5.110150730361098	merxlebeAND
27	401	Q0	LA082389-0142	27	5.096443390896712	merxlebeAND
28	401	Q0	LA010489-0026	28	5.086599738681089	merxlebeAND
29	401	Q0	LA021590-0208	29	5.070719867194082	merxlebeAND
30	401	Q0	LA022290-0221	30	5.0381508547565685	merxlebeAND
31	401	Q0	LA100289-0097	31	5.037974546489284	merxlebeAND
32	401	Q0	LA041290-0088	32	5.025387408972495	merxlebeAND
33	401	Q0	LA030990-0070	33	5.021008826321708	merxlebeAND
34	401	Q0	LA040689-0200	34	5.016233306183249	merxlebeAND
35	401	Q0	LA030690-0168	35	5.015236799714735	merxlebeAND
36	401	Q0	LA120489-0110	36	5.011677817493935	merxlebeAND
37	401	Q0	LA121589-0087	37	5.003802403664816	merxlebeAND
38	401	Q0	LA100189-0189	38	4.997092130780651	merxlebeAND
39	401	Q0	LA020690-0078	39	4.996780802615824	merxlebeAND
40	401	Q0	LA011289-0071	40	4.994797474781549	merxlebeAND
41	401	Q0	LA031490-0158	41	4.993503766204197	merxlebeAND
42	401	Q0	LA050889-0081	42	4.9926448470116656	merxlebeAND
43	401	Q0	LA120689-0034	43	4.988344867167488	merxlebeAND
44	401	Q0	LA111089-0142	44	4.972261265372243	merxlebeAND
45	401	Q0	LA100890-0076	45	4.96563088015223	merxlebeAND
46	401	Q0	LA100190-0033	46	4.959097760451025	merxlebeAND
47	401	Q0	LA030790-0048	47	4.955181895485336	merxlebeAND
48	401	Q0	LA061089-0056	48	4.94860218599951	merxlebeAND
49	401	Q0	LA083090-0247	49	4.9421797218951316	merxlebeAND
50	401	Q0	LA111289-0073	50	4.935030586177118	merxlebeAND
51	401	Q0	LA091289-0149	51	4.931766167483419	merxlebeAND
52	401	Q0	LA101990-0026	52	4.9232339472981	merxlebeAND
53	401	Q0	LA071790-0033	53	4.912806381325571	merxlebeAND
...

Topic:	401 AP:	0.0468 P@10:	0.1 NDCG@10:	0.0948 NDCG@1000:	0.3455
Topic:	402 AP:	0.0607 P@10:	0.2 NDCG@10:	0.2935 NDCG@1000:	0.2688
Topic:	403 AP:	0.5075 P@10:	0.6 NDCG@10:	0.5302 NDCG@1000:	0.7407
Topic:	404 AP:	0.004 P@10:	0 NDCG@10:	0 NDCG@1000:	0.144
Topic:	405 AP:	0.0347 P@10:	0.2 NDCG@10:	0.1518 NDCG@1000:	0.1014
Topic:	406 AP:	0.3941 P@10:	0.2 NDCG@10:	0.3833 NDCG@1000:	0.7362
Topic:	407 AP:	0.2083 P@10:	0.6 NDCG@10:	0.8424 NDCG@1000:	0.5775
Topic:	408 AP:	0.0941 P@10:	0.4 NDCG@10:	0.5369 NDCG@1000:	0.383
Topic:	409 AP:	0.1 P@10:	0.1 NDCG@10:	0.2891 NDCG@1000:	0.2891
Topic:	410 AP:	1 P@10:	0.4 NDCG@10:	1 NDCG@1000:	1
Topic:	411 AP:	0.0983 P@10:	0.2 NDCG@10:	0.3052 NDCG@1000:	0.3759
Topic:	412 AP:	0.3809 P@10:	0.7 NDCG@10:	0.85 NDCG@1000:	0.6379
Topic:	413 AP:	0.0101 P@10:	0 NDCG@10:	0 NDCG@1000:	0.1505
Topic:	414 AP:	0.092 P@10:	0.1 NDCG@10:	0.1872 NDCG@1000:	0.3251
Topic:	415 AP:	0.25 P@10:	0.1 NDCG@10:	0.3904 NDCG@1000:	0.3904
Topic:	417 AP:	0.339 P@10:	0.7 NDCG@10:	0.7885 NDCG@1000:	0.7172
Topic:	418 AP:	0.1304 P@10:	0.6 NDCG@10:	0.7223 NDCG@1000:	0.3433
Topic:	419 AP:	0.5084 P@10:	0.2 NDCG@10:	0.6307 NDCG@1000:	0.7353
Topic:	420 AP:	0.623 P@10:	0.8 NDCG@10:	0.8522 NDCG@1000:	0.8976
Topic:	421 AP:	0.018 P@10:	0 NDCG@10:	0 NDCG@1000:	0.2446
Topic:	422 AP:	0.3624 P@10:	0.4 NDCG@10:	0.3251 NDCG@1000:	0.6452
Topic:	424 AP:	0.0206 P@10:	0.2 NDCG@10:	0.1428 NDCG@1000:	0.1155
Topic:	425 AP:	0.2851 P@10:	0.6 NDCG@10:	0.6316 NDCG@1000:	0.5957
Topic:	426 AP:	0.0274 P@10:	0.3 NDCG@10:	0.2221 NDCG@1000:	0.1068
Topic:	427 AP:	0.0707 P@10:	0.2 NDCG@10:	0.3052 NDCG@1000:	0.3041
Topic:	428 AP:	0.2533 P@10:	0.1 NDCG@10:	0.3904 NDCG@1000:	0.4821
Topic:	429 AP:	0.2827 P@10:	0.1 NDCG@10:	0.3904 NDCG@1000:	0.5748
Topic:	430 AP:	0.4811 P@10:	0.4 NDCG@10:	0.6278 NDCG@1000:	0.6823
Topic:	431 AP:	0.0984 P@10:	0 NDCG@10:	0 NDCG@1000:	0.3387
Topic:	432 AP:	0.0036 P@10:	0 NDCG@10:	0 NDCG@1000:	0.1034
Topic:	433 AP:	0.0052 P@10:	0 NDCG@10:	0 NDCG@1000:	0.1109
Topic:	434 AP:	0.85 P@10:	0.1 NDCG@10:	0.6131 NDCG@1000:	0.7527
Topic:	435 AP:	0.064 P@10:	0.1 NDCG@10:	0.0784 NDCG@1000:	0.3038
Topic:	436 AP:	0.0370 P@10:	0.3 NDCG@10:	0.2528 NDCG@1000:	0.203
Topic:	438 AP:	0.0954 P@10:	0.2 NDCG@10:	0.1357 NDCG@1000:	0.4320
Topic:	439 AP:	0.0038 P@10:	0 NDCG@10:	0 NDCG@1000:	0.1264
Topic:	440 AP:	0.5823 P@10:	0.6 NDCG@10:	0.6937 NDCG@1000:	0.8198
Topic:	441 AP:	0.8079 P@10:	0.6 NDCG@10:	0.7381 NDCG@1000:	0.7381
Topic:	442 AP:	0.0252 P@10:	0.2 NDCG@10:	0.1447 NDCG@1000:	0.2129
Topic:	443 AP:	0.1059 P@10:	0.2 NDCG@10:	0.2173 NDCG@1000:	0.4526
Topic:	445 AP:	0.2444 P@10:	0.2 NDCG@10:	0.4162 NDCG@1000:	0.4162
Topic:	446 AP:	0.0292 P@10:	0 NDCG@10:	0 NDCG@1000:	0.2083
Topic:	448 AP:	0.0175 P@10:	0 NDCG@10:	0 NDCG@1000:	0.2111
Topic:	449 AP:	0.0095 P@10:	0 NDCG@10:	0 NDCG@1000:	0.0099
Topic:	450 AP:	0.2235 P@10:	0.4 NDCG@10:	0.3031 NDCG@1000:	0.5881

7c)

The results of the baseline BM25 retrieval can be found at hw4-bm25-stem-merxlebe.txt

results >		hw4-bm25-stem-merxlebe.txt			
1	401 Q0	LA030690-0168	1	6.571483823576605	merxlebeAND
2	401 Q0	LA050690-0109	2	6.460391036324291	merxlebeAND
3	401 Q0	LA041290-0088	3	6.195006855170041	merxlebeAND
4	401 Q0	LA040989-0072	4	6.172135707298894	merxlebeAND
5	401 Q0	LA110690-0075	5	5.705613351575752	merxlebeAND
6	401 Q0	LA050789-0068	6	5.661913089819003	merxlebeAND
7	401 Q0	LA100890-0131	7	5.579819203949144	merxlebeAND
8	401 Q0	LA021490-0049	8	5.571769295008746	merxlebeAND
9	401 Q0	LA031590-0102	9	5.52464248254968	merxlebeAND
10	401 Q0	LA050790-0042	10	5.51258405708243	merxlebeAND
11	401 Q0	LA071890-0073	11	5.39108117431041	merxlebeAND
12	401 Q0	LA070989-0066	12	5.377856459418849	merxlebeAND
13	401 Q0	LA040590-0157	13	5.333143411638125	merxlebeAND
14	401 Q0	LA111089-0188	14	5.326337522258674	merxlebeAND
15	401 Q0	LA021890-0100	15	5.3156148123297715	merxlebeAND
16	401 Q0	LA021190-0168	16	5.260124850054082	merxlebeAND
17	401 Q0	LA060890-0011	17	5.248969539016584	merxlebeAND
18	401 Q0	LA030990-0189	18	5.228508594961867	merxlebeAND
19	401 Q0	LA112990-0229	19	5.221609137526233	merxlebeAND
20	401 Q0	LA050990-0043	20	5.220355308790698	merxlebeAND
21	401 Q0	LA062789-0019	21	5.218694409415905	merxlebeAND
22	401 Q0	LA061290-0074	22	5.216610031634419	merxlebeAND
23	401 Q0	LA102390-0068	23	5.201782239675228	merxlebeAND
24	401 Q0	LA072590-0049	24	5.201028527677608	merxlebeAND
25	401 Q0	LA091490-0080	25	5.152900208458081	merxlebeAND
26	401 Q0	LA031490-0158	26	5.146041020270798	merxlebeAND
27	401 Q0	LA062290-0172	27	5.1296766207072615	merxlebeAND
28	401 Q0	LA100490-0081	28	5.128459397144992	merxlebeAND
29	401 Q0	LA072790-0172	29	5.1232437985758885	merxlebeAND
30	401 Q0	LA051190-0084	30	5.123202529715591	merxlebeAND
31	401 Q0	LA100289-0097	31	5.0999981800200604	merxlebeAND
32	401 Q0	LA060489-0038	32	5.09475966080801	merxlebeAND
33	401 Q0	LA050390-0176	33	5.091258099252266	merxlebeAND
34	401 Q0	LA043090-0043	34	5.091009908575692	merxlebeAND
35	401 Q0	LA031389-0076	35	5.077286244562838	merxlebeAND
36	401 Q0	LA021590-0209	36	5.068460595146931	merxlebeAND
37	401 Q0	LA082389-0142	37	5.05189632226523	merxlebeAND
38	401 Q0	LA010489-0026	38	5.043525314263908	merxlebeAND
39	401 Q0	LA021490-0144	39	5.043488003913742	merxlebeAND
40	401 Q0	LA021590-0208	40	5.027858461604249	merxlebeAND
41	401 Q0	LA022790-0091	41	5.006450673232987	merxlebeAND
42	401 Q0	LA122190-0057	42	5.00254208025757	merxlebeAND
43	401 Q0	LA121789-0230	43	4.99618587779891	merxlebeAND
44	401 Q0	LA050889-0081	44	4.994612483680944	merxlebeAND
45	401 Q0	LA121589-0087	45	4.9889579658041034	merxlebeAND
46	401 Q0	LA022290-0221	46	4.987029355657999	merxlebeAND
47	401 Q0	LA030990-0070	47	4.979825651025676	merxlebeAND
48	401 Q0	LA011990-0099	48	4.974834551970854	merxlebeAND
49	401 Q0	LA120489-0110	49	4.9682016182297035	merxlebeAND
50	401 Q0	LA110990-0013	50	4.96766878028147	merxlebeAND
51	401 Q0	LA042490-0209	51	4.96231996063267	merxlebeAND
52	401 Q0	LA040689-0200	52	4.958702081620088	merxlebeAND
53	401 Q0	LA100189-0189	53	4.957842202393272	merxlebeAND
54	401 Q0	LA020690-0078	54	4.955621916806248	merxlebeAND

Topic:	401 AP:	0.1032 P@10:	0.3 NDCG@10:	0.2468 NDCG@1000:	0.45
Topic:	402 AP:	0.2070 P@10:	0.3 NDCG@10:	0.3838 NDCG@1000:	0.6017
Topic:	403 AP:	0.5075 P@10:	0.6 NDCG@10:	0.5302 NDCG@1000:	0.7407
Topic:	404 AP:	0.0097 P@10:	0 NDCG@10:	0 NDCG@1000:	0.1604
Topic:	405 AP:	0.0284 P@10:	0.1 NDCG@10:	0.0734 NDCG@1000:	0.1476
Topic:	406 AP:	0.4470 P@10:	0.3 NDCG@10:	0.4817 NDCG@1000:	0.77
Topic:	407 AP:	0.1800 P@10:	0.4 NDCG@10:	0.5038 NDCG@1000:	0.5224
Topic:	408 AP:	0.130 P@10:	0.3 NDCG@10:	0.3571 NDCG@1000:	0.4019
Topic:	409 AP:	0.1 P@10:	0.1 NDCG@10:	0.2891 NDCG@1000:	0.2891
Topic:	410 AP:	1 P@10:	0.4 NDCG@10:	1 NDCG@1000:	1
Topic:	411 AP:	0.1700 P@10:	0.3 NDCG@10:	0.4441 NDCG@1000:	0.4782
Topic:	412 AP:	0.4480 P@10:	0.6 NDCG@10:	0.7183 NDCG@1000:	0.7271
Topic:	413 AP:	0.0833 P@10:	0 NDCG@10:	0 NDCG@1000:	0.2702
Topic:	414 AP:	0.1055 P@10:	0.1 NDCG@10:	0.2021 NDCG@1000:	0.344
Topic:	415 AP:	0.25 P@10:	0.1 NDCG@10:	0.3904 NDCG@1000:	0.3904
Topic:	417 AP:	0.3500 P@10:	0.7 NDCG@10:	0.7792 NDCG@1000:	0.7333
Topic:	418 AP:	0.2863 P@10:	0.6 NDCG@10:	0.7273 NDCG@1000:	0.622
Topic:	419 AP:	0.5633 P@10:	0.3 NDCG@10:	0.7542 NDCG@1000:	0.7542
Topic:	420 AP:	0.82 P@10:	0.8 NDCG@10:	0.8904 NDCG@1000:	0.8974
Topic:	421 AP:	0.0180 P@10:	0 NDCG@10:	0 NDCG@1000:	0.2887
Topic:	422 AP:	0.3750 P@10:	0.7 NDCG@10:	0.6559 NDCG@1000:	0.6621
Topic:	424 AP:	0.1544 P@10:	0.1 NDCG@10:	0.0638 NDCG@1000:	0.593
Topic:	425 AP:	0.4815 P@10:	0.7 NDCG@10:	0.7824 NDCG@1000:	0.8230
Topic:	426 AP:	0.0342 P@10:	0.1 NDCG@10:	0.0784 NDCG@1000:	0.1708
Topic:	427 AP:	0.0072 P@10:	0.2 NDCG@10:	0.3052 NDCG@1000:	0.3838
Topic:	428 AP:	0.1094 P@10:	0.1 NDCG@10:	0.1092 NDCG@1000:	0.3354
Topic:	429 AP:	0.7080 P@10:	0.4 NDCG@10:	0.9223 NDCG@1000:	0.9223
Topic:	430 AP:	0.8203 P@10:	0.4 NDCG@10:	0.8608 NDCG@1000:	0.7499
Topic:	431 AP:	0.3195 P@10:	0.6 NDCG@10:	0.8084 NDCG@1000:	0.6536
Topic:	432 AP:	0.0017 P@10:	0 NDCG@10:	0 NDCG@1000:	0.066
Topic:	433 AP:	0.0051 P@10:	0 NDCG@10:	0 NDCG@1000:	0.1102
Topic:	434 AP:	0.5417 P@10:	0.1 NDCG@10:	0.6131 NDCG@1000:	0.7452
Topic:	435 AP:	0.0411 P@10:	0 NDCG@10:	0 NDCG@1000:	0.2813
Topic:	436 AP:	0.0925 P@10:	0.7 NDCG@10:	0.7191 NDCG@1000:	0.3182
Topic:	438 AP:	0.1158 P@10:	0.1 NDCG@10:	0.0851 NDCG@1000:	0.4570
Topic:	439 AP:	0.0140 P@10:	0 NDCG@10:	0 NDCG@1000:	0.1801
Topic:	440 AP:	0.5803 P@10:	0.5 NDCG@10:	0.6274 NDCG@1000:	0.8131
Topic:	441 AP:	0.8400 P@10:	0.6 NDCG@10:	0.7834 NDCG@1000:	0.7834
Topic:	442 AP:	0.0231 P@10:	0.2 NDCG@10:	0.137 NDCG@1000:	0.1976
Topic:	443 AP:	0.1054 P@10:	0.2 NDCG@10:	0.1763 NDCG@1000:	0.4203
Topic:	445 AP:	0.2444 P@10:	0.2 NDCG@10:	0.4162 NDCG@1000:	0.4162
Topic:	446 AP:	0.0245 P@10:	0 NDCG@10:	0 NDCG@1000:	0.2323
Topic:	448 AP:	0.0062 P@10:	0 NDCG@10:	0 NDCG@1000:	0.1947
Topic:	449 AP:	0.0075 P@10:	0 NDCG@10:	0 NDCG@1000:	0.0913
Topic:	450 AP:	0.2393 P@10:	0.4 NDCG@10:	0.4381 NDCG@1000:	0.6497

7d)

After obtaining the BM25 results for the baseline and stem versions, I calculated the four effectiveness measures to evaluate the two results. This can be seen in the following table:

Run	Mean AP	Mean P@10	Mean NDCG@10	Mean NDCG@1000
baseline	0.209	0.253	0.335	0.425
stem	0.251	0.284	0.374	0.487
p-value	0.005167	0.150811	0.100977	0.001687

As we can see, the Porter Stemmer version of the BM25 results has a higher Mean Average Precision, Mean Precision@10, Mean NDCG@10 and Mean NDCG@1000. This shows an overall improvement over the baseline run. I utilized the t-test to compare the statistical significance of the difference between these two results. As we can see, the p-values computed are extremely small for each effectiveness measure. However, if we are to apply the $p < 0.05$ concept to these results, we can see that only Mean Average Precision and Mean NDCG@1000 are statistically significant. Although the other two evaluation metrics are close to 0.05, this test shows that the different ins not statistically significant. This shows that the Porter Stemmer

version is better in overall Mean Average Precision and Mean NDCG@1000. As we can see, these two evaluation metrics have something in common. Both evaluation metrics are evaluating a large sample size (Mean Average Precision is across the entire results set, and NDCG @ 1000 is evaluating at rank 1000), versus the other metrics evaluating a smaller sample size (both at rank 10). This demonstrates that the Porter Stemmer's better results become more statistically significant the bigger sample size is.

To see how on a topic-by-topic basis BM25 is in the baseline and stemmed results, I compared each topic's mean effectiveness measure. I utilized Excel to highlight if the two are equivalent, then the column will equal 0 and highlight yellow. If stemmed results are better than the baseline results, then the column equals 1 and highlights green. If stemmed results are worse than the baseline results, then the column equals -1 and highlights red.

Baseline					Stem					Comparison				
Topic	Mean AP	Mean P@10	Mean NDCG@1	Mean NDCG@10	Topic	Mean AP	Mean P@10	Mean NDCG@1	Mean NDCG@10	Mean AP	Mean P@10	Mean NDCG@1	Mean NDCG@1000	
401	0.0468	0.1	0.0948	0.3455	401	0.1032	0.3	0.2486	0.45	1	1	1	1	
402	0.0607	0.2	0.2935	0.2668	402	0.2078	0.3	0.3836	0.6017	1	1	1	1	
403	0.5075	0.6	0.5302	0.7407	403	0.5075	0.6	0.5302	0.7407	0	0	0	0	
404	0.004	0	0	0.144	404	0.0097	0	0	0.1694	1	0	0	1	
405	0.0347	0.2	0.1518	0.1614	405	0.0264	0.1	0.0734	0.1476	-1	-1	-1	-1	
406	0.3941	0.2	0.3833	0.7382	406	0.4479	0.3	0.4617	0.77	1	1	1	1	
407	0.2063	0.6	0.6434	0.5775	407	0.1689	0.4	0.5036	0.5224	-1	-1	-1	-1	
408	0.0941	0.4	0.5389	0.363	408	0.136	0.3	0.3571	0.4619	1	-1	-1	1	
409	0.1	0.1	0.2891	0.2891	409	0.1	0.1	0.2891	0.2891	0	0	0	0	
410	1	0.4	1	1	410	1	0.4	1	1	0	0	0	0	
411	0.0983	0.2	0.3052	0.3759	411	0.1799	0.3	0.4441	0.4762	1	1	1	1	
412	0.3609	0.7	0.65	0.6379	412	0.4468	0.8	0.7163	0.7271	1	1	1	1	
413	0.0101	0	0	0.1505	413	0.0833	0	0	0.2702	1	0	0	1	
414	0.092	0.1	0.1672	0.3251	414	0.1055	0.1	0.2021	0.344	1	0	1	1	
415	0.25	0.1	0.3904	0.3904	415	0.25	0.1	0.3904	0.3904	0	0	0	0	
417	0.339	0.7	0.7885	0.7172	417	0.3569	0.7	0.7792	0.7333	1	0	-1	1	
418	0.1394	0.6	0.7223	0.3433	418	0.2663	0.6	0.7273	0.622	1	0	1	1	
419	0.5084	0.2	0.6367	0.7353	419	0.5833	0.3	0.7542	0.7542	1	1	1	1	
420	0.623	0.8	0.8522	0.8876	420	0.62	0.8	0.8604	0.8874	-1	0	1	-1	
421	0.018	0	0	0.2446	421	0.0189	0	0	0.2857	1	0	0	1	
422	0.3624	0.4	0.3251	0.6452	422	0.3756	0.7	0.5559	0.6621	1	1	1	1	
424	0.0206	0.2	0.1428	0.1155	424	0.1544	0.1	0.0636	0.593	1	-1	-1	1	
425	0.2851	0.6	0.6316	0.5657	425	0.4818	0.7	0.7624	0.8239	1	1	1	1	
426	0.0274	0.3	0.2221	0.1696	426	0.0342	0.1	0.0784	0.1798	1	-1	-1	1	
427	0.0797	0.2	0.3052	0.3041	427	0.0972	0.2	0.3052	0.3838	1	0	0	1	
428	0.2533	0.1	0.3904	0.4821	428	0.1094	0.1	0.1952	0.3354	-1	0	-1	-1	
429	0.2827	0.1	0.3904	0.5748	429	0.7986	0.4	0.9223	0.9223	1	1	1	1	
430	0.4811	0.4	0.6278	0.6823	430	0.6203	0.4	0.6608	0.7499	1	0	1	1	
431	0.0984	0	0	0.3367	431	0.3195	0.6	0.6064	0.6838	1	1	1	1	
432	0.0036	0	0	0.1034	432	0.0017	0	0	0.068	-1	0	0	-1	
433	0.0052	0	0	0.1109	433	0.0051	0	0	0.1102	-1	0	0	-1	
434	0.55	0.1	0.6131	0.7527	434	0.5417	0.1	0.6131	0.7452	-1	0	0	-1	
435	0.064	0.1	0.0784	0.3638	435	0.0411	0	0	0.2813	-1	-1	-1	-1	
436	0.0376	0.3	0.2528	0.203	436	0.0925	0.7	0.7191	0.3182	1	1	1	1	
438	0.0954	0.2	0.1357	0.4326	438	0.1158	0.1	0.0851	0.4579	1	-1	-1	1	
439	0.0038	0	0	0.1264	439	0.0146	0	0	0.1681	1	0	0	1	
440	0.5823	0.6	0.6937	0.8198	440	0.5883	0.5	0.6274	0.8131	-1	-1	-1	-1	
441	0.6079	0.6	0.7381	0.7381	441	0.6496	0.6	0.7634	0.7634	1	0	1	1	
442	0.0252	0.2	0.1447	0.2129	442	0.0231	0.2	0.137	0.1976	-1	0	-1	-1	
443	0.1089	0.2	0.2173	0.4326	443	0.1054	0.2	0.1763	0.4203	-1	0	-1	-1	
445	0.2444	0.2	0.4162	0.4162	445	0.2444	0.2	0.4162	0.4162	0	0	0	0	
446	0.0292	0	0	0.2063	446	0.0245	0	0	0.2323	-1	0	0	1	
448	0.0175	0	0	0.2111	448	0.0092	0	0	0.1947	-1	0	0	-1	
449	0.0095	0	0	0.0969	449	0.0075	0	0	0.0913	-1	0	0	-1	
450	0.2235	0.4	0.3031	0.5861	450	0.2393	0.4	0.4381	0.6497	1	0	1	1	

As we can see, topics: 405, 407, 420, 428, 432, 433, 434, 435, 440, 442, 443, 446, 448, 449 are overall equal or better in the baseline than the stemmed results. When looking at these queries, it seems like queries using plurals and -ing endings are more common amongst these results.

This shows that when the user makes a search with these additions to their words, it may be more likely that the results lead to non-relevant documents being returned.